

Estimating Unobserved COVID-19 Infections in the United States

Quinn White

Submitted to the Department of Statistical and Data Sciences
of Smith College
in partial fulfillment
of the requirements for the degree of
Bachelor of Arts

Ben Baumer, Primary Faculty Advisor
Nicholas Reich, Secondary Faculty Advisor

May 2023

Acknowledgements

Will add

Table of Contents

Chapter 1: Motivation	1
Chapter 2: Background	3
2.1 Probabalistic Bias Analysis	3
2.2 Background for the Approach	4
2.2.1 Simple Discrete Example	5
2.2.2 General Solution for the Discrete Case	7
2.2.3 General Solution for the Continuous Case	7
2.2.4 Implementation through the Sampling-Importance-Resampling Algorithm	8
2.3 Bayesian Melding Applied to COVID-19 Misclassification	9
2.3.1 Distribution of $\theta = \{\alpha, \beta, P(S_1 \text{untested})\}$	10
2.3.2 Direct Prior and Induced Prior Distributions for $P(S_0 \text{test}_+, \text{untested})$	10
2.3.3 Pooling	11
2.3.4 Derivation of M	13
2.4 Sampling-Importance-Resampling Algorithm	15
2.4.1 Overview	15
2.4.2 Proof that Algorithm Obtains Approximate Sample from Target Distribution	18
2.4.3 Obtaining Logarithmic Pooled Distribution with the Sampling-Importance-Resampling Algorithm	19
2.5 LOESS Smoothing	21
2.5.1 Introduction	21
2.5.2 Fitting the LOESS Curve	22
2.6 Kernel Density Estimation	26
Chapter 3: Definition of Prior Distributions for Bias Parameters	27
3.1 Background on the Beta Distribution	28
3.2 Background on the Gamma Distribution	28
3.3 Definition of Prior Distributions for Incomplete Testing Correction	28
3.3.1 Defining $P(S_1 Untested)$	28
3.3.2 Defining α	28
3.3.3 Defining β	28
3.3.4 Defining $P(S_0 test_+, untested)$	28
3.4 Definition of Priors for Test Inaccuracy Correction	28

3.4.1	Defining Test Sensitivity (S_e)	28
3.5	Defining Test Specificity (S_p)	28
3.6	Summary Table of Bias Parameter Distributions	28
3.7	Correction for Incomplete Testing	28
3.8	Correction for Diagnostic Test Inaccuracy	28
3.8.1	Derivation of Formula for Correction for Diagnostic Test In- accuracy	28
Chapter 4:	Details of Implementation	29
Chapter 5:	Results	31
5.1	County-level	31
5.2	State-level	31
5.3	Comparison to the Covidestim Model	31
5.3.1	Overview	31
5.3.2	The Covidestim Model	31
5.3.3	Assumptions	31
5.4	Comparison to Other Indicators	31
5.5	Seropositivity Data	31
5.6	County-level	31
5.7	State-level	31
Appendix A:	Appendix	33
A.1	Derivation of the Mean and Variance of the Beta Distribution	33
References	35

Abstract

As we have navigated the COVID-19 pandemic, case counts have been a central source of information for understanding transmission dynamics and the effect of public health interventions. However, because the number of cases we observe is limited by the testing effort in a given location, the case counts presented on local or national dashboards are only a fraction of the true infections. Variations in testing rate by time and location impacts the number of cases that go unobserved, which can cloud our understanding of the true COVID-19 incidence at a given time point and can create biases in downstream analyses. Additionally, the number of cases we observe is impacted by the sensitivity and specificity of the diagnostic test. To quantify the number of true infections given incomplete testing and diagnostic test inaccuracy, this work implements probabilistic bias analysis at a biweekly time scale from January 1, 2021 through February 2022. In doing so, we can estimate a range of possible true infections for a given time interval and location. This approach can be applied at the state level across the United States, as well as in some counties where the needed data are available.

Dedication

You can have a dedication here if you wish.

Chapter 1

Motivation

Placeholder

Chapter 2

Background

2.1 Probabilistic Bias Analysis

Often the focus of quantifying error about an effect estimate focuses on random error rather than the systematic error. For example, typical frequentist confidence intervals are frequent in medical and epidemiological literature, although they have faced rising criticism (Greenland et al., 2016). These confidence intervals quantify the fraction of the times we expect the true value to fall in this interval under the assumption that our model is correct. That is, if we ran an experiment 100 times and computed the effect size each time, we would expect the 95% confidence interval to contain the true value to 95 of those times, on average. Neyman stressed this in his original publication formalizing the concept of a confidence interval in 1937 (Neyman, 1937). The nuance that the confidence interval is not the probability that the true value falls within this interval, however, is often lost in the discussion of results, in part because the true meaning of a confidence interval is less intuitive.

The aim of quantitative bias analysis is to estimate systematic error to give a range of possible values for the true quantity of interest. In this sense, it is a type of sensitivity analysis. It can be used to estimate various kinds of biases, from misclassification, as is implemented in this work, as well as selection bias and unmeasured confounding (Petersen, Ranker, Barnard-Mayers, MacLehose, & Fox, 2021). Often, the goal of performing such an analysis is to see how these sources of bias affect our estimates; in particular, under what situations of bias the observed effect would be null.

There are multiple different forms of bias analysis (Lash, Fox, & Fink, 2009). The most simple case, simple bias analysis, is correcting a point estimate for a single source of error. Multidimensional bias analysis extends this to consider sets of bias parameters, but still provides a corrected point estimate rather than a range of plausible estimates. Probabilistic bias analysis, meanwhile, defines probability distributions for bias parameters to generate a distribution of corrected estimates by repeatedly correcting estimates for bias under different combinations of the parameter values. Then, via Monte Carlo we obtain a distribution of corrected estimates that reflect the corrected values under different scenarios of bias, that is, under dif-

ferent combinations of the bias parameters. This can give us a better idea for the extent of uncertainty about the corrected estimates, although this uncertainty does depend on the specification of the bias parameter distributions. Inherent in bias analysis is the dependence of our results on the specification of bias parameters, which reflect what is known from available data, literature, or theory on the extent of bias that may occur. There is uncertainty about how we define these distributions or values; otherwise, if the precise values of the bias parameters were known, we could simply correct the estimates and probabilistic bias analysis would not be useful.

Although some forms of probabilistic bias analysis can be applied to summarized data, for example, frequencies in a contingency table, the methods are most often implemented with unsummarized data in its original form, as implemented here.

In choosing specific distributions for the bias parameters, different specifications may yield density functions where most of the density is within a similar interval, which means the choice of the specific distribution will not be sensitive to the particular choice of density.

2.2 Background for the Approach

The Bayesian melding approach was proposed by Poole et al. (Poole & Raftery, 2000).

This approach enables us to account for both uncertainty from inputs and outputs of a deterministic model. The initial motivation for the approach was to study the population dynamics of whales in the presence of substantial uncertainty around model inputs for population growth (Poole & Raftery, 2000). However, the framework provided by Poole et al. can be applied in any circumstance where we have uncertainty around some quantities θ and ϕ where there is a deterministic function $M : \theta \rightarrow \phi$. Due to the utility of Bayesian melding in various contexts, since this deterministic model M could take on a wide range of forms, the approach has since been applied in various fields, including urban simulations (Ševčíková, Raftery, & Waddell, 2007), ecology (Robson, 2014), and infectious disease (Powers et al., 2011).

Let $M : \theta \rightarrow \phi$ be the deterministic model defined by the function relating a vector of input parameters θ to an output vector ϕ , and suppose we have a prior on θ denoted $f_\theta(\theta)$ and a prior on ϕ denoted $f_\phi^{direct}(\phi)$.

However, note that we actually have two distinct priors on ϕ . There is the prior formed by the distribution induced on ϕ by the prior for θ and the function M , where we denote this induced prior $f_\phi^{induced}(\phi)$. Generally, these priors are based on different sources of information.

If M^{-1} exists, we can write this induced prior $f_\phi^{induced}(\phi) = f_\theta(M^{-1}(\phi))|J(\phi)|^1$. This result follows from the fact $M(\theta) = \phi$, so we apply a change of variables to

¹In the continuous case we need to multiply by $|J(\phi)|$, but not in the discrete case (Blitzstein & Hwang, 2019).

obtain the distribution of ϕ from the distribution of $M(\theta)$.

In practice, M^{-1} rarely exists since θ is often of higher dimensionality than ϕ , in which cases M is not invertible. This means we generally approximate f_ϕ^{induced} without acquiring its analytical form.

In addition to this induced prior, we have the prior $f_\phi^{\text{direct}}(\phi)$, which does not involve M nor the inputs θ . Since these priors are based on different sources of information and may reflect different uncertainties, often it is useful to use both sources of information to inform our estimates. To do so, we need to combine the distributions for f_ϕ^{induced} and f_ϕ^{direct} to create a pooled distribution.

Multiple pooling strategies exist for distinct distributions, but one requirement for a Bayesian analysis is that the distribution should be independent of the order in which the prior is updated and the combining of the prior distribution. That is, updating the prior distributions using Bayes' theorem and then combining distributions should yield the same result as combining distributions and then updating this combined distribution; pooling methods that have this property are deemed externally Bayesian. Logarithmic pooling has been shown to be externally Bayesian under some conditions, which are likely to hold in most settings. Furthermore, logarithmic pooling has actually been shown to be the only pooling method where this holds (Genest, McConway, & Schervish, 1986). For this reason, Poole *et al.* recommend proceeding with logarithmic pooling for Bayesian melding.

The logarithmically pooled prior for ϕ by pooling f_ϕ^{induced} and f_ϕ^{direct} is


$$f_\phi^{\text{pooled}}(\phi) = t(\alpha)(f_\phi^{\text{induced}}(\phi))^\alpha(f_\phi^{\text{direct}}(\phi))^{1-\alpha}$$

where $\alpha \in [0, 1]$ is a pooling weight and $t(\alpha)$ is the normalizing constant. Commonly, a choice of $\alpha = 0.5$ is used to give the priors equal weight. In this case, logarithmic pooling may be referred to as geometric pooling since it is equivalent to taking a geometric mean.

If M is invertible, we can obtain the constrained distributions for the model inputs by simply inverting M . However, M is rarely invertible, so we have to think about how to proceed in the noninvertible case.

2.2.1 Simple Discrete Example

To get intuition for a valid strategy Poole *et al.* recommend, we consider a mapping $M : \theta \rightarrow \phi$ for $\theta \in \mathbb{R}$ and $\phi \in \mathbb{R}$ defined as follows (Figure 2.1). Note the choice of $f_\theta, f_\phi^{\text{direct}}$ does not matter here as long as they are valid densities.



θ	$f_\theta(\theta)$	$M(\theta) = \phi$	$f_\phi^{\text{direct}}(\phi)$
1	0.3	1	0.4
2	0.2	2	0.6
3	0.5	2	0.6

Figure 2.1: A simple discrete example where M is not invertible.

We see that M is not invertible since $\theta = 1$ and $\theta = 2$ both map to $\phi = 2$, which implies the inverse M^{-1} would not be well defined.

We can generate a sample from the density f_ϕ^{induced} by sampling from f_θ and computing $M(\theta)$.

So we have

$$\begin{aligned} f_\phi^{\text{induced}}(1) &= f_\theta(1) = 0.3 && (\text{since } \theta = 1 \text{ maps } \phi = 1) \\ f_\phi^{\text{induced}}(2) &= f_\theta(2) + f_\theta(3) = 0.2 + 0.5 = 0.7 && (\text{since } \theta = 2 \text{ and } \theta = 3 \text{ both map to } \phi = 2) \end{aligned}$$

Then, we can compute the logarithmically pooled prior with $\alpha = 0.5$ by taking $f_\phi^{\text{induced}}(\phi)^\alpha f_\phi^{\text{direct}}(\phi)^{1-\alpha}$.

This gives us

$$\begin{aligned} f_\phi^{\text{induced}}(\phi)^\alpha f_\phi^{\text{direct}}(\phi)^{1-\alpha} &= (0.3)^{0.5}(0.4)^{0.5} = 0.3464 \\ f_\phi^{\text{induced}}(\phi)^\alpha f_\phi^{\text{direct}}(\phi)^{1-\alpha} &= (0.7)^{0.5}(0.6)^{0.5} = 0.6481. \end{aligned}$$

To make this a valid density, however, these probabilities must sum to 1, so we renormalize by dividing by $(0.3464 + 0.6481)$. Denoting the pooled prior in phi-space as $f_\phi^{\text{pooled}}(\phi)$, this gives us

$$\begin{aligned} f_\phi^{\text{pooled}}(1) &= \frac{0.3464}{0.3464 + 0.6481} = 0.3483 \\ f_\phi^{\text{pooled}}(2) &= \frac{0.6481}{0.3464 + 0.6481} = 0.6517. \end{aligned}$$

We summarize these results and compare f_ϕ^{induced} , f_ϕ^{direct} , and f_ϕ^{pooled} in Figure 2.2.




Figure 2.2

However, we also want the pooled prior on the inputs θ , that is, $f_\theta^{\text{pooled}}(\theta)$.

Poole et al. reasoned as follows. Since M uniquely maps $\theta = 1$ to $\phi = 1$, the probability that $\theta = 1$ should be equal to the probability $\phi = 1$. That is, we should have $f_\theta^{\text{pooled}}(1) = f_\phi^{\text{pooled}}(1)$.

However, the relationship for $\theta = 2$ or $\theta = 3$ to ϕ is not one to one. Since $M(2) = 2$ and $M(3) = 2$, the sum of the probabilities for $\theta = 1$ and $\theta = 2$ should be equal to that for $\phi = 2$, that is, $f_\theta^{\text{pooled}}(2) + f_\theta^{\text{pooled}}(3) = f_\phi^{\text{pooled}}(2) = 0.6517$.

The challenge here is how we divide the probability for $f_\phi^{pooled}(2)$, which is defined, among $f_\theta^{pooled}(2)$ and $f_\theta^{pooled}(3)$. The prior for ϕ yields no information to assist in this choice, because knowing which value ϕ takes on does not give us any information about whether $\theta = 2$ or $\theta = 3$. Thus, the information we have about θ must be taken from $f_\theta(\theta)$.

That is, we can assign a probability for $f_\theta^{pooled}(2)$ by considering the probability that $\theta = 2$ relative to the probability $\theta = 3$, computing

$$f_\theta^{pooled}(2) = f_\phi^{pooled}(2) \left(\frac{f_\theta(2)}{f_\theta(2) + f_\theta(3)} \right).$$

That is, if the probability θ takes on the value 2 is lower in this case than the probability $\theta = 3$ which we know from the prior on θ , $f_\theta(\theta)$, then the pooled prior on θ , $f_\theta^{pooled}(2)$, should reflect this.

Using this reasoning, we have

$$\begin{aligned} f_\theta^{pooled}(2) &= (0.7) \frac{0.2}{0.2 + 0.5} = 0.1862 \\ f_\theta^{pooled}(3) &= (0.7) \frac{0.5}{0.2 + 0.5} = 0.4655. \end{aligned}$$

The result in this simple example, using $f_\theta(\theta)$ to determine how to distribute the probability for values of ϕ where multiple θ map to ϕ , can be used to derive general formulas to compute $f_\theta^{pooled}(\theta)$ for discrete and continuous distributions (Poole & Raftery, 2000).

2.2.2 General Solution for the Discrete Case

Denote the possible values of θ as A_1, A_2, \dots , the possible values of ϕ as B_1, B_2, \dots , and a mapping $m : \mathbb{N} \rightarrow \mathbb{N}$ such that $M(A_i) = B_{m(i)}$ and $C_j = M^{-1}(B_j) = \{A_i : M(A_i) = B_j\}$. Then

$$f_\theta^{pooled}(A_i) = f_\phi^{pooled}(B_{m(i)}) \left(\frac{f_\theta(A_i)}{f_\phi^{induced}(B_{m(i)})} \right).$$

2.2.3 General Solution for the Continuous Case

We denote $B = M(A) = \{M(\theta) : \theta \in A\}$ and $C = M^{-1}(B) = \{\theta : M(\theta) \in B\}$.

Then

$$f_\phi^{pooled}(M(\theta)) = t(\alpha) f_\theta(\theta) \left(\frac{f_\phi^{direct}(M(\theta))}{f_\phi^{induced}(M(\theta))} \right)^{1-\alpha} \quad (2)$$

where $t(\alpha)$ is a renormalizing constant for the choice of α .

2.2.4 Implementation through the Sampling-Importance-Resampling Algorithm

We can obtain the pooled distributions f_{θ}^{pooled} and f_{ϕ}^{pooled} by using the Sampling-Importance-Resampling Algorithm.

The steps are as follows.

1. We draw θ from its prior distribution $f_{\theta}(\theta)$.
2. For every θ_i we compute $\phi_i = M(\theta_i)$ to obtain a sample from the induced distribution.
3. Since the density $f_{\phi}^{induced}(\phi)$ is unlikely to have an analytical form, we can compute it via a density approximation such as kernel density estimation.
4. Construct weights proportional to the ratio of the prior on ϕ evaluated at $M(\theta_i)$ to the induced prior $f_{\phi}^{induced}$ evaluated at $M(\theta_i)$. If a likelihood $L_1(\theta)$ for the inputs and $L_2(\phi)$ is available, the weights are

$$w_i = \left(\frac{f_{\phi}^{direct}(M(\theta_i))}{f_{\phi}^{induced}(M(\theta_i))} \right)^{1-\alpha}.$$

However, in this work, no likelihood is available for the variables of interest, so the likelihood is left out of the weights, leaving us with

$$w_i = \left(\frac{f_{\phi}^{direct}(M(\theta_i))}{f_{\phi}^{induced}(M(\theta_i))} \right)^{1-\alpha}.$$

5. Sample θ and ϕ from step (1) with probabilities proportional to the weights from (4).

2.3 Bayesian Melding Applied to COVID-19 Misclassification

In this work, we can relate the inputs $\theta = \{P(S_1|\text{untested}), \alpha, \beta\}$ and $\phi = P(S_0|\text{test}_+, \text{untested})$ by the deterministic model $M : \theta \rightarrow \phi$ given by $P(S_0|\text{test}_+, \text{untested}) = \frac{\beta(1 - P(S_1|\text{untested}))}{\beta(1 - P(S_1|\text{untested})) + \alpha P(S_1|\text{untested})}$. The derivation of M is in the [following section](#).

Now, we have two distributions on ϕ : the distribution based on data on the asymptomatic rate of infection of COVID-19, and the distribution formed by taking $M(\theta)$ where θ represents the values from the defined distributions of α, β , and $P(S_1|\text{untested})$. With Bayesian melding, we pool these distributions using logarithmic pooling, and then implement the sampling-importance-resampling algorithm to obtain constrained distributions of the inputs θ that are in accordance with information about the asymptomatic rate of the virus.

Due to the uncertainty around our definitions of α and β , it is particularly useful to leverage the information we have about the asymptomatic rate of the virus $P(S_0|\text{test}_+, \text{untested})$ because a large collection of studies has been published in this area. In a meta-analysis pooling data from 95 studies, the pooled estimate among the confirmed population that was asymptomatic was 40.50% [95% CI, 33.50%-47.50%] (Ma et al., 2021). Another meta-analysis including 350 studies estimated the asymptomatic percentage to be 36.9% [95% CI: 31.8 to 42.4%], and, when restricting to screening studies, 47.3% (95% CI: 34.0% -61.0%) (Sah et al., 2021).

This means we have two priors on the asymptomatic rate ϕ , that by taking $M(\theta)$ for sampled values of θ , denoted f_ϕ^{induced} in the previous section, and that based on data about the asymptomatic rate, f_ϕ^{direct} .

2.3.1 Distribution of $\theta = \{\alpha, \beta, P(S_1|\text{untested})\}$

First, we obtain a sample $\theta_1, \theta_2, \dots, \theta_k$ from θ (Figure 2.3).

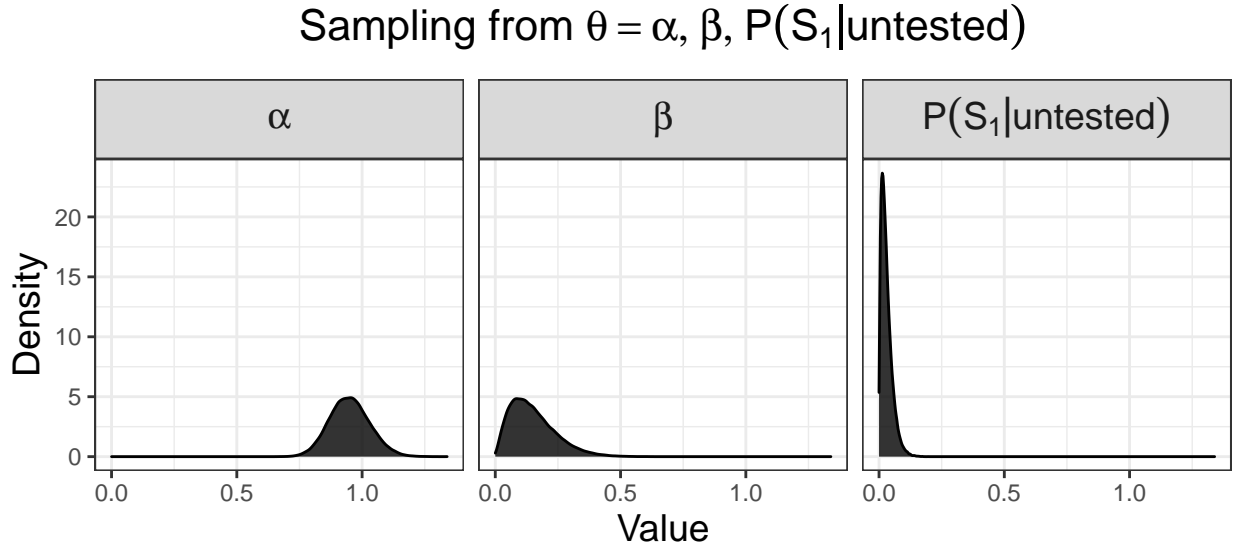


Figure 2.3

2.3.2 Direct Prior and Induced Prior Distributions for $P(S_0|\text{test}_+, \text{untested})$

Then, taking $M(\theta)$, we can compute the induced distribution $f_\phi^{\text{induced}}(M(\theta))$ and compare it to our prior on ϕ from meta-analyses on the asymptomatic rate, $f_\phi^{\text{direct}}(\phi)$ (2.4).




Figure 2.4

2.3.3 Pooling

At this point, we want to obtain the logarithmically pooled distribution, denoted f_{ϕ}^{pooled} .

Now, as described in greater detail in the section on the [Sampling-Importance-Resampling algorithm](#), the weights are $w_i = \left(\frac{f_{\phi}^{\text{direct}}(M(\theta_i))}{f_{\phi}^{\text{induced}}(M(\theta_i))} \right)^{1-\alpha}$.

We perform a kernel density estimation to approximate the density of $f_{\phi}^{\text{induced}}(\phi)$ at the coordinates ϕ_1, \dots, ϕ_M . To compute $f_{\phi}^{\text{direct}}(\phi)$, we can use the density function f_{ϕ}^{direct} .

Once we have these weights, we resample the ϕ_1, \dots, ϕ_M to obtain a sample from the target distribution $t(\alpha) \left(f^{\text{induced}}(M(\theta)) \right)^{0.5} \left(f^{\text{direct}}(M(\theta)) \right)^{0.5}$, where $t(\alpha)$ is the normalizing constant needed to make the pooled density valid. We resample $\theta_1, \dots, \theta_k$ with the same weights to obtain the constrained distributions for the inputs.

We see the melded distributions and pre-melding distributions in Figure 2.5.

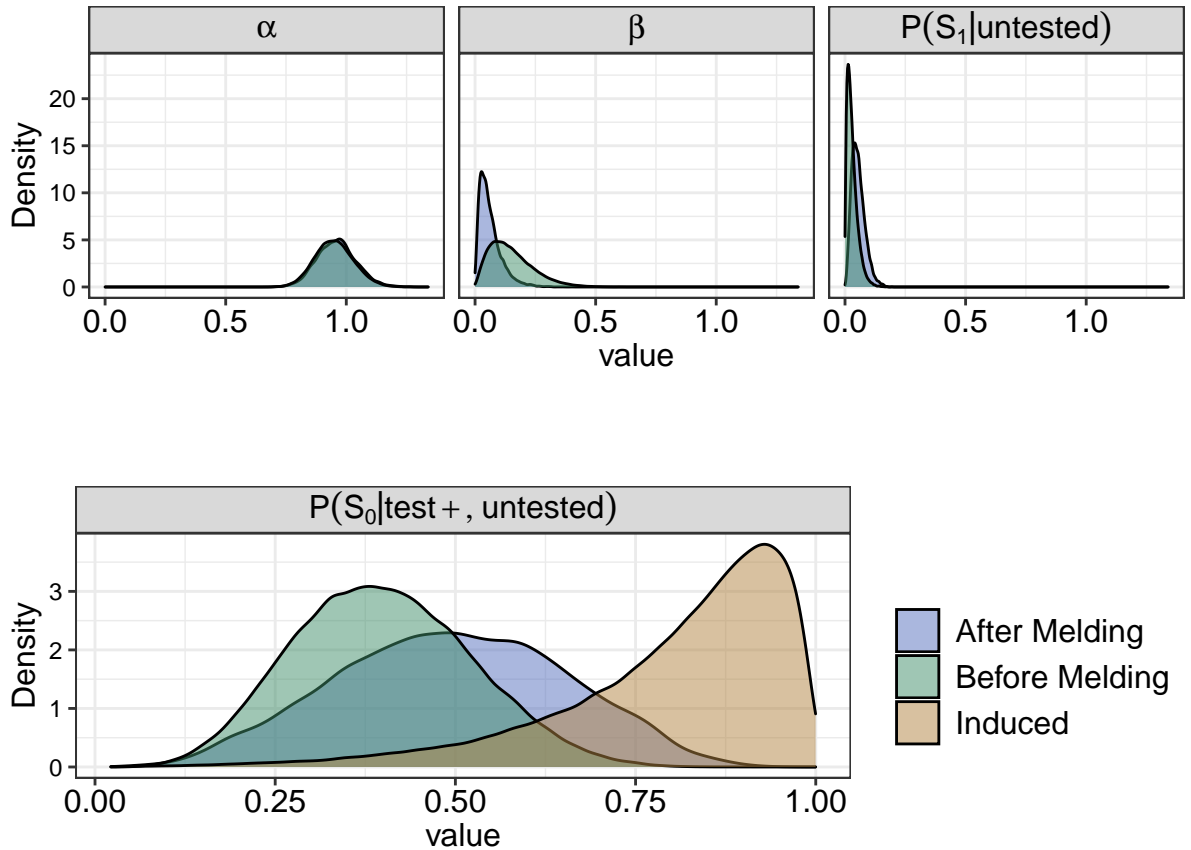


Figure 2.5

Comparing the induced and direct priors on $P(S_0|\text{test}_+, \text{untested})$ above, we see that although they have shared support, some values from the induced distribution we acquire by using M to generate values of ϕ from sampled values of θ are very unlikely to be in accordance with the information we know about the prevalence of SARS-CoV-2 asymptomatic infection. This is where Bayesian melding comes into play. Pooling these distributions enable us to take both the prior on f^{direct} from published analyses on asymptomatic infection, and the induced prior, f^{induced} , into account to constrain the distributions of both the model inputs $\theta = \{\alpha, \beta, P(S_1|\text{untested})\}$ and model output $\phi = P(S_0|\text{test}_+, \text{untested})$ to be in accordance with both prior distributions. We then use these constrained distributions as inputs in the probabilistic bias analysis.

2.3.4 Derivation of M

We define θ as the set of bias parameters $\{P(S_1|\text{untested}), \alpha, \beta\}$. The parameters α and β relate the observed overall test positivity rate to the test positivity rate we would obtain if we tested the asymptomatic and symptomatic partitions of the untested population. We define:

- $\alpha = \frac{P(\text{test}_+|S_1, \text{untested})}{P(\text{test}_+|\text{tested})}$
- $\beta = \frac{P(\text{test}_+|S_0, \text{untested})}{P(\text{test}_+|\text{tested})}$.

The parameter $P(S_1|\text{untested})$ reflects the probability someone among the untested population has moderate to severe COVID-like symptoms.

We relate this set of parameters to the asymptomatic infection rate $\phi = P(S_0|\text{test}_+, \text{untested})$ by the function $M : \theta \rightarrow \phi$:

$$M(\theta) = \frac{\beta(1 - P(S_1|\text{untested}))}{\beta(1 - P(S_1|\text{untested})) + \alpha(P(S_1|\text{untested}))} = P(S_0|\text{test}_+, \text{untested}).$$

In what follows, we show this equality holds.

Since we have $\alpha = \frac{P(\text{test}_+|S_1, \text{untested})}{P(\text{test}_+|\text{tested})}$ and $\beta = \frac{P(\text{test}_+|S_0, \text{untested})}{P(\text{test}_+|\text{tested})}$, we can write

$$\begin{aligned} &= \frac{\frac{P(\text{test}_+|S_0, \text{untested})}{P(\text{test}_+|\text{tested})}(1 - P(S_1|\text{untested}))}{\frac{P(\text{test}_+|S_0, \text{untested})}{P(\text{test}_+|\text{tested})}(1 - P(S_1|\text{untested})) + \frac{P(\text{test}_+|S_1, \text{untested})}{P(\text{test}_+|\text{tested})}P(S_1|\text{untested})} \end{aligned}$$

and cancelling out the term $P(\text{test}_+|\text{tested})$ we have

$$= \frac{P(\text{test}_+|S_0, \text{untested})(1 - P(S_1|\text{untested}))}{P(\text{test}_+|S_0, \text{untested})(1 - P(S_1|\text{untested})) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}.$$

Since $P(S_0|\text{untested}) = 1 - P(S_1|\text{untested})$,

$$= \frac{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested})}{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}.$$

Applying the definition of conditional probability to the term $P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested})$

in the numerator,

$$\begin{aligned}
& \left(\frac{P(\text{test}_+, S_0, \text{untested})}{P(S_0, \text{untested})} \right) \left(\frac{P(S_0, \text{untested})}{P(\text{untested})} \right) \\
= & \frac{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}{\left(\frac{P(\text{test}_+, S_0, \text{untested})}{P(S_0, \text{untested})} \right) \left(\frac{P(S_0, \text{untested})}{P(\text{untested})} \right)} \\
= & \frac{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}{\frac{P(\text{test}_+, S_0, \text{untested})}{P(\text{untested})}} \\
= & \frac{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}{P(\text{test}_+, S_0|\text{untested})} \\
= & \frac{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}{P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}.
\end{aligned}$$

We can substitute this result in for the $P(\text{test}_+|S_0, \text{untested})P(S_0|\text{untested})$ term in the denominator to yield

$$= \frac{P(\text{test}_+, S_0|\text{untested})}{P(\text{test}_+, S_0|\text{untested}) + P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested})}$$

With same reasoning, we can simplify

$$P(\text{test}_+|S_1, \text{untested})P(S_1|\text{untested}) = P(S_1, \text{test}_+|\text{untested}),$$

giving us

$$\begin{aligned}
& \frac{P(\text{test}_+, S_0|\text{untested})}{P(\text{test}_+, S_0|\text{untested}) + P(S_1, \text{test}_+|\text{untested})} \\
= & \frac{P(\text{test}_+, S_0|\text{untested})}{P(\text{test}_+|\text{untested})} \\
& \frac{P(S_0, \text{test}_+, \text{untested})}{P(\text{untested})} \\
= & \frac{P(\text{test}_+, \text{untested})}{P(\text{untested})} \\
& \frac{P(S_0, \text{test}_+, \text{untested})}{P(\text{test}_+, \text{untested})} \\
= & P(S_0|\text{test}_+, \text{untested}).
\end{aligned}$$

Hence, we have

$$P(S_0|\text{test}_+, \text{untested}) = \frac{\beta(1 - P(S_1|\text{untested}))}{\beta(1 - P(S_1|\text{untested})) + \alpha(P(S_1|\text{untested}))}$$

as desired. \square

2.4 Sampling-Importance-Resampling Algorithm

2.4.1 Overview

The Sampling-Importance-Resampling Algorithm, introduced in Rubin (1987), is a non-iterative method for approximating a sample from a target probability density function f .

The two main steps are the sampling step and importance resampling step. We have two (generally distinct) sample sizes, where m is the initial sample size and r is the posterior sample size.

In the sampling step, we draw an independent and identically distributed sample of size m from g , Y_1, Y_2, \dots, Y_m . Then, we compute weights $h(Y)$ such that $g \cdot h \propto f$. That is, we set the weights

$$w_i = h(Y_i) = \frac{\frac{f(Y_i)}{g(Y_i)}}{\sum_{i=1}^m \frac{f(Y_i)}{g(Y_i)}}.$$

We resample with these defined weights to obtain a sample of size r from Y_1, Y_2, \dots, Y_m . We denote this resample Z_1, \dots, Z_r .

The method is most efficient when g is a good approximation of f . Also of interest is setting the sampling size and posterior sample size. The algorithm generates independent and identically distributed samples as $m/r \rightarrow \infty$, but in most applications m/r between 10 and 20 is appropriate (Rubin, Gelman, & Meng, 2004).

Sometimes the resulting distribution will have a closed form, which we can see in the following two examples.

Example 1:

Suppose $Y \sim \text{Exp}(\lambda)$, so we have the PDF $f_Y(y) = \lambda e^{-\lambda y}$, and we sample Z_1, \dots, Z_r from Y_1, \dots, Y_m with weights direction proportional to X , that is, $h(Y) = Y$.

Then Z_1, \dots, Z_r is approximately a sample from $h(x) f_Y(y) = y \lambda e^{-\lambda y}$.

From the PDF of the gamma distribution, $\frac{\beta^\alpha}{\Gamma(\alpha)} y^{\alpha-1} e^{-\beta y}$ we can recognize that $y \cdot e^{-\lambda y}$ corresponds to the gamma distribution with $\alpha = 2$ and $\beta = \lambda$.

We can see this result by considering Y before and after resampling below (Figure 2.6).

Sampling from $Y \sim \text{Exp}(\lambda = 0.2)$ with Weights $h(y) = y$




Figure 2.6

Then, we can see that the PDF of the the gamma distribution with $\alpha = 2$ and $\beta = \lambda$ corresponds to the post-sampling distribution as expected (Figure 2.7).

Sampling from $Y \sim \text{Exp}(\lambda = 0.2)$ with Weights $h(y) = y$

PDF of Gamma(2, 0.2) in Red




Figure 2.7

Example 2:

Similarly, again suppose $Y \sim \text{Exp}(\lambda)$, so $f_Y(y) = \lambda e^{-\lambda y}$. However, now we sample with weights defined by $h(y) = e^{-\lambda y}$. Then our sample Z_1, \dots, Z_r is approximately a sample from

$$\begin{aligned} h(y) f_Y(y) &= e^{-\lambda y} \cdot \lambda e^{-\lambda y} \\ &= e^{-2\lambda y} \end{aligned}$$

which is proportional to the exponential distribution with parameter 2λ .

The distributions before and after resampling are shown in Figure 2.8.

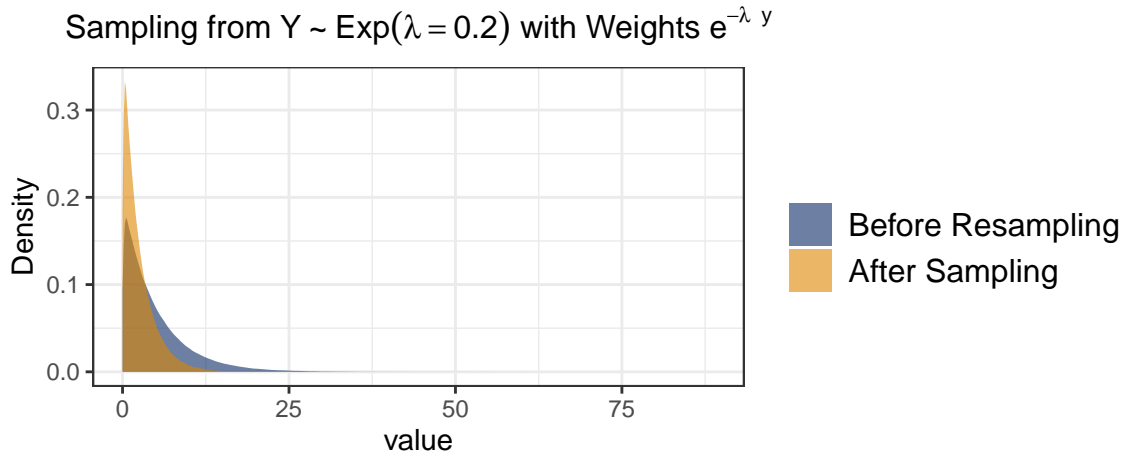


Figure 2.8

and then plotting the PDF of the exponential distribution with parameter 2λ we can see the correspondence to the post-sampling distribution (Figure 2.9).




Figure 2.9

2.4.2 Proof that Algorithm Obtains Approximate Sample from Target Distribution

The Sampling-Importance-Resampling is fundamental to the implementation of Bayesian melding. To gain further insight into how sampling with weights $w_i = \left(\frac{f_\phi^{\text{direct}}(M(\theta_i))}{f_\phi^{\text{induced}}(M(\theta_i))} \right)^{0.5}$ approximates a sample from the target distribution the logarithmically pooled distribution f^{pooled} , we first prove a more general result.

Suppose we sample Y_1, Y_2, \dots, Y_m independently and identically distributed with probability density function g and compute the weights

$$w_i = \frac{h(Y_i)}{\sum_{i=1}^m h(Y_i)}$$

for some nonnegative function h defined on the support of Y .

If we sample Z_1, \dots, Z_r from the discrete distribution Y_1, \dots, Y_m such that

$$P(Z = Y_i) = \frac{h(Y_i)}{\sum_{i=1}^m h(Y_i)} = w_i,$$

then Z_1, \dots, Z_r is approximately a sample with density proportional to $h \cdot g$.

Since Z is sampled from Y , we have

$$P(Z \leq x) = \sum_{z_i \leq x} P(Z = z_i) = \sum_{Y_i \leq x} P(Z = Y_i).$$

We can take this sum to be over all possible values of Y by including the indicator function $\mathbb{I}(Y_i \leq x)$, yielding

$$= \sum_{i=1}^m P(Z = y_i) \mathbb{I}(Y_i \leq x).$$

and since $P(Z = Y_i) = \frac{h(Y_i)}{\sum_{i=1}^m h(Y_i)}$ by definition we have

$$\begin{aligned} &= \sum_{i=1}^m \frac{h(Y_i)}{\sum_{i=1}^m h(Y_i)} \mathbb{I}(Y_i \leq x) \\ &= \left(\frac{1}{\sum_{i=1}^m h(Y_i)} \right) \sum_{i=1}^m h(Y_i) \mathbb{I}(Y_i \leq x) \\ &= \frac{\sum_{i=1}^m h(Y_i) \mathbb{I}(Y_i \leq x)}{\sum_{i=1}^m h(Y_i)} \\ &= \frac{\frac{1}{m} \sum_{i=1}^m h(Y_i) \mathbb{I}(Y_i \leq x)}{\frac{1}{m} \sum_{i=1}^m h(Y_i)}. \end{aligned}$$

Now, we need the Weak Law of Large Numbers. That is, if we have a sequence of random variables X_1, X_2, \dots with finite variance, then,

$$\lim_{n \rightarrow \infty} \left(\frac{1}{n} \sum_{i=1}^n X_i \right) = E(X_i).$$

Applying this law to both the numerator and denominator, we obtain

$$\begin{aligned} \lim_{m \rightarrow \infty} \left(\frac{\frac{1}{m} \sum_{i=1}^m h(Y_i) \mathbb{I}(Y_i \leq x)}{\frac{1}{m} \sum_{i=1}^m h(Y_i)} \right) &= \frac{E_g[h(Y) \mathbb{I}(Y \leq x)]}{E_g[h(Y)]} \\ &= \frac{\int_{-\infty}^{\infty} h(y) \mathbb{I}(y \leq x) g(y) dy}{\int_{-\infty}^{\infty} h(y) g(y) dy} \\ &= \frac{\int_{-\infty}^x h(y) g(y) dy}{\int_{-\infty}^{\infty} h(y) g(y) dy} \\ &\propto \int_{-\infty}^x h(y) g(y) dy. \end{aligned}$$

It follows that the probability density function of Z is proportional to $h \cdot g$.

□

2.4.3 Obtaining Logarithmic Pooled Distribution with the Sampling-Importance-Resampling Algorithm

As outlined in Carvalho, Villela, Coelho, & Bastos (2023), we can formally define logarithmic pooling as follows.

If we have a set of densities $\{f_1(\phi), f_2(\phi), \dots, f_n(\phi)\}$ and corresponding pooling weights $\alpha = \{\alpha_1, \alpha_2, \dots, \alpha_n\}$, then the pooled density is

$$t(\alpha) \prod_{i=1}^n f_i(\phi)^{\alpha_i}$$

where $t(\alpha)$ is the normalizing constant $t(\alpha) = \frac{1}{\int_{\Phi} \prod_{i=1}^n f_i(\phi)^{\alpha_i} d\phi}$ to ensure the pooled density is a valid probability density.

The case for this work is more simple: we only have two densities we wish to pool, $f_{\phi}^{\text{induced}}$ and f_{ϕ}^{direct} , and we assign them equal weights by letting $\alpha = \{.5, .5\}$. This yields

$$f^{\text{pooled}}(\phi) = t(\alpha) \left(f^{\text{induced}}(\phi) \right)^{0.5} \left(f^{\text{direct}}(\phi) \right)^{0.5}.$$

Since our target distribution is $t(\alpha) \left(f^{\text{induced}}(\phi) \right)^{0.5} \left(f^{\text{direct}}(\phi) \right)^{0.5}$, and we have a

sample from $f^{induced}$, we compute the weights such that

$$\begin{aligned} w_i &\propto \frac{\left(f^{induced}(\phi_i)\right)^{0.5} \left(f^{direct}(\phi_i)\right)^{0.5}}{f^{induced}(\phi_i)} \\ &= \frac{\left(f^{direct}(\phi_i)\right)^{0.5}}{\left(f^{induced}(\phi_i)\right)^{0.5}} \\ &= \left(\frac{f^{direct}(\phi_i)}{f^{induced}(\phi_i)}\right)^{0.5}. \end{aligned}$$

Sampling from $f^{induced}$ with these weights will yield a sample with approximately the target density $t(\alpha) \left(f^{induced}(\phi)\right)^{0.5} \left(f^{direct}(\phi)\right)^{0.5}$ from the result in the [previous section](#).

2.5 LOESS Smoothing

2.5.1 Introduction

Locally estimated scatterplot smoothing (LOESS) fits a collection of local regression models to obtain a smooth curve through the observed data (Figure 2.10). It is highly flexible in the sense that we do not have to specify the functional relationship between the predictor and response variable for the entire range of the predictor, which may be impossible in various settings. It is particularly useful when working with time series data with substantial noise.




Figure 2.10

To perform LOESS smoothing, we estimate a set of local regressions (Chambers, 1997). To do this, we must specify the span; this smoothing parameter is the fraction of the data that is used for the local polynomial fit. With a smaller span, the resulting curve will fit the trends more closely, while a larger span reflects broader trends (Figure 2.11).

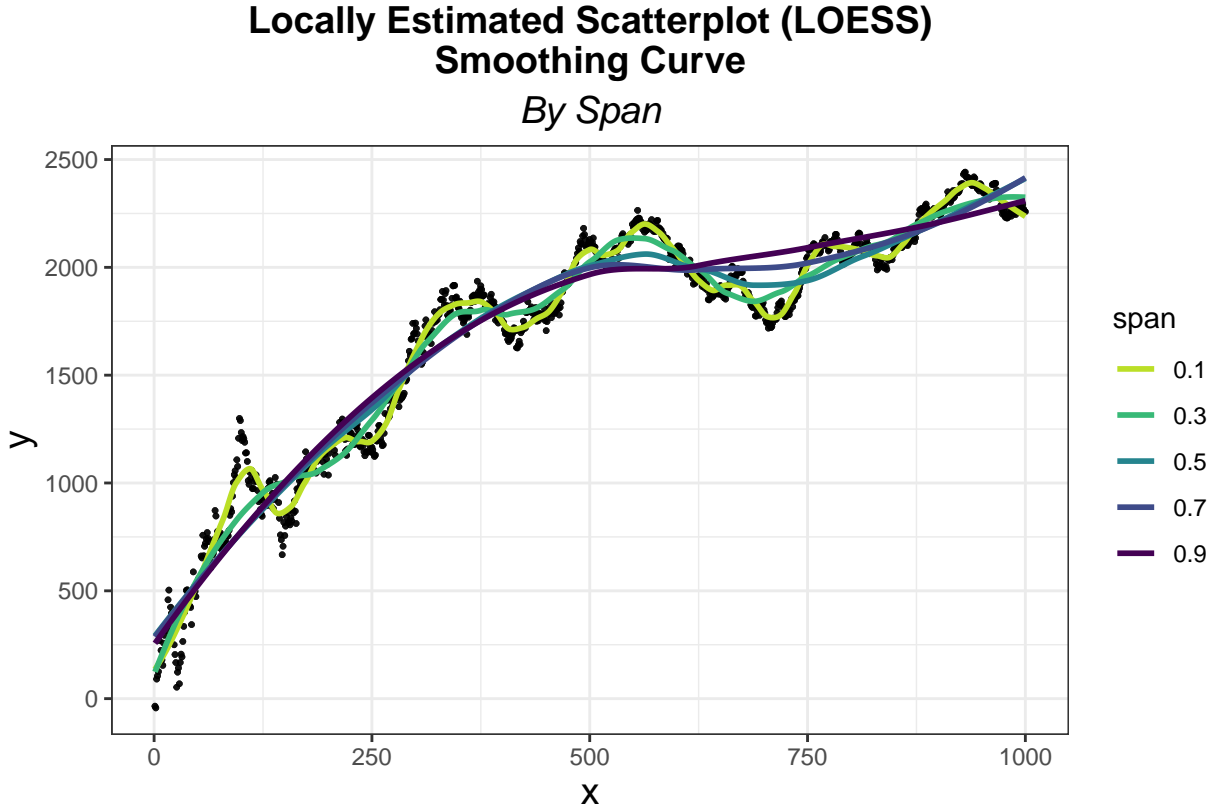


Figure 2.11

2.5.2 Fitting the LOESS Curve

To introduce some notation for the model at hand, we have a dependent variable y and independent variable x , where y and x are related by some unknown function g , that is, $y = g(x) + \epsilon^2$. When we want to use LOESS smoothing to estimate g , often this function is complex, so we break up the problem into estimating a set of local regressions.

To obtain a predicted value $\hat{g}(x^*)$ for a particular value of the independent variable x^* , we fit a polynomial with greatest weight placed on points in the neighborhood of x^* , where the width of this neighborhood is defined by the choice of smoothing span. Let $\alpha \in (0, 1]$ denote the chosen smoothing span.

For a particular value of x^* , we estimate the predicted value $\hat{g}(x^*)$ by fitting a local regression. We first compute the weights by computing the vector of distances from this point x^* , that is,

$$\Delta(x^*) = |\mathbf{x} - x^*|$$

We define $q = \text{floor}(\alpha n)$, and take $\Delta_q(x^*) \in \mathbb{R}$ to be the q^{th} smallest distance of $\Delta(x^*)$.

²Recall we use bold type for vectors, e.g., $\mathbf{x} \in \mathbb{R}^n$ is a vector with observations $x_i \in \mathbb{R}$.

The vector of weights is then

$$T(\Delta(x^*), \Delta_q(x^*))$$

where T is the tricube weight function given by

$$T(x) = \begin{cases} (1 - (x)^3)^3 & \text{for } |x| < 1 \\ 0 & \text{for } |x| \geq 1 \end{cases}.$$

Essentially, this process gives weight to points in the neighborhood of x^* . Consider $x^* = 500$ and smoothing span = $\alpha = .2$.

Then the weights we obtain are given in Figure 2.12.

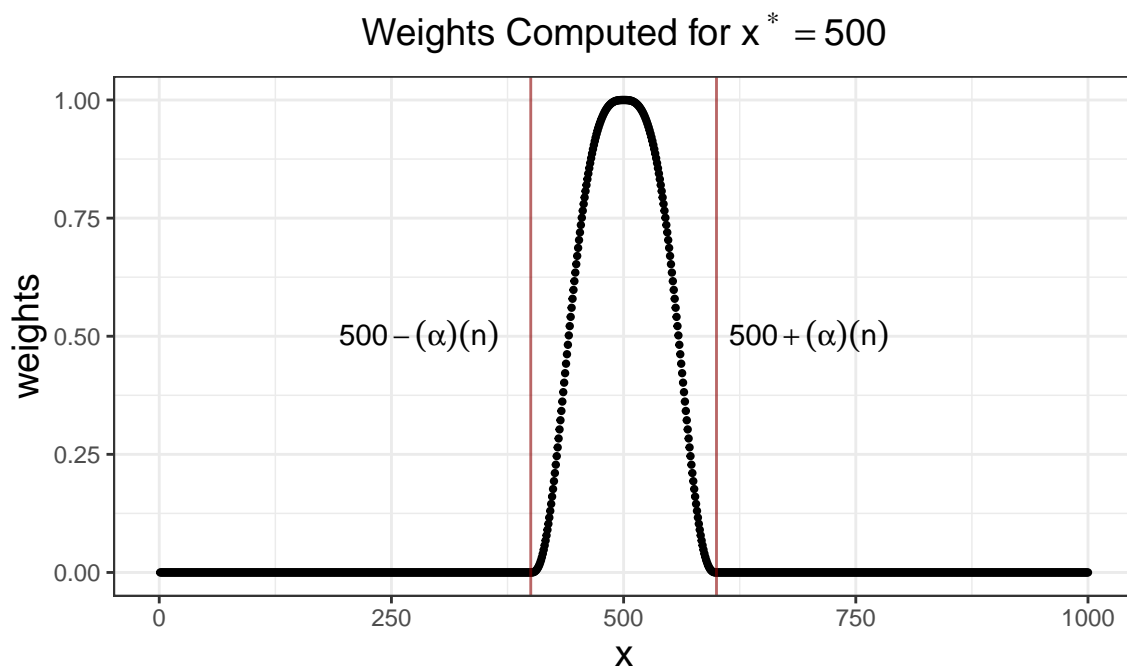


Figure 2.12: The only values with nonzero weights are those within the interval $(500 - \alpha(n), 500 + \alpha(n))$. That is, the proportion α of the data points closest to x^* will have nonzero weights.

We fit a linear regression with polynomial terms, typically with degree up to 2, with these weights. For example, fitting the model for this same $x^* = 500$, we obtain the polynomial in Figure 2.13.

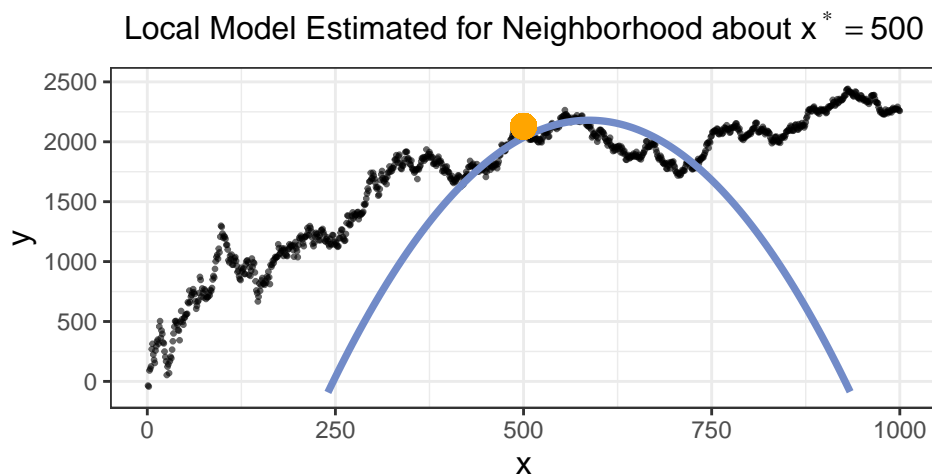


Figure 2.13

By fitting the model for every point in x , we obtain the smoothed line shown in red in Figure 2.14.




Figure 2.14

Smoothing methods are sensitive to the choice of smoothing parameter h , which represents the fraction of the data that is used for the local polynomial fit.

Methods exist for picking the smoothing parameter h that minimizes the mean squared error between the predicted values from the estimated line and observed values of the dependent variable, for example, leave-one-out cross-validation or generalized cross-validation.

However, for this work, we used LOESS smoothing to smooth survey data from the COVID-19 Trends and Impact Survey (Reinhart et al., 2021). We choose the smoothing parameter for each variable based on domain knowledge regarding the level of noise present for each variable of interest. For example, there is substantial noise in the screening test positivity data that reflect trends that do not represent meaningful differences in the screening test positivity. Some trends in the screening sensitivity may be due to scheduled workplace screenings happening at regular time intervals, and some of the variation may be due to the frequency of screening testing due to other variables, such as the access and cost of testing.

Since the ratio $\frac{\text{screening test positivity}}{\text{overall test positivity}}$ is used to estimate $\beta = \frac{P(\text{test}_+|S_0, \text{untested})}{P(\text{test}_+|\text{tested})}$, the variability in the screening positivity creates substantial variability in our estimates of β .

In light of this variability and the presence of other trends regarding the screening test positivity, we set the span to $\frac{4}{12} = 0.33$ to fit the local regressions for 4-month intervals with the aim to capture the broader trends over time.

INCLUDE FIGURE OF SMOOTHED ESTIMATES HERE

There was less variability in the smoothing span for the weighted percentage of COVID-like illness, the estimate of $P(S_1|\text{untested})$. Hence, we set the smoothing parameter to 0.2 to detect trends at a finer time scale.

Sensitivity analyses with modified versions of the smoothing span of β are included in the appendix in the section INCLUDE SECTION.

2.6 Kernel Density Estimation

Chapter 3

Definition of Prior Distributions for Bias Parameters

Placeholder

3.1 Background on the Beta Distribution

3.2 Background on the Gamma Distribution

3.3 Definition of Prior Distributions for Incomplete Testing Correction

3.3.1 Defining $P(S_1|Untested)$

3.3.2 Defining α

3.3.3 Defining β

3.3.4 Defining $P(S_0|test+, untested)$

3.4 Definition of Priors for Test Inaccuracy Correction

3.4.1 Defining Test Sensitivity (S_e)

3.5 Defining Test Specificity (S_p)

3.6 Summary Table of Bias Parameter Distributions

3.7 Correction for Incomplete Testing

3.8 Correction for Diagnostic Test Inaccuracy

3.8.1 Derivation of Formula for Correction for Diagnostic Test Inaccuracy

Chapter 4

Details of Implementation

- Describe each step here
- Mention reproducible workflow with make

Chapter 5

Results

Placeholder

5.1 County-level

5.2 State-level

5.3 Comparison to the Covidestim Model

5.3.1 Overview

5.3.2 The Covidestim Model

5.3.3 Assumptions

5.4 Comparison to Other Indicators

5.5 Seropositivity Data

5.6 County-level

5.7 State-level

Appendix A

Appendix

A.1 Derivation of the Mean and Variance of the Beta Distribution

To add

References

Placeholder

- Blitzstein, J. K., & Hwang, J. (2019). *Introduction to probability* (Second edition). Boca Raton: CRC Press.
- Carvalho, L. M., Villela, D. A. M., Coelho, F. C., & Bastos, L. S. (2023). Bayesian Inference for the Weights in Logarithmic Pooling. *Bayesian Analysis*, 18(1). <http://doi.org/10.1214/22-BA1311>
- Chambers, J. M. (Ed.). (1997). *Statistical models in S* (Reprint). London: Chapman & Hall.
- Genest, C., McConway, K. J., & Schervish, M. J. (1986). Characterization of Externally Bayesian Pooling Operators. *The Annals of Statistics*, 14(2), 487–501. Retrieved from <https://www.jstor.org/stable/2241231>
- Greenland, S., Senn, S. J., Rothman, K. J., Carlin, J. B., Poole, C., Goodman, S. N., & Altman, D. G. (2016). Statistical tests, P values, confidence intervals, and power: a guide to misinterpretations. *European Journal of Epidemiology*, 31(4), 337–350. <http://doi.org/10.1007/s10654-016-0149-3>
- Lash, T. L., Fox, M. P., & Fink, A. K. (2009). *Applying Quantitative Bias Analysis to Epidemiologic Data*. New York, NY: Springer New York. <http://doi.org/10.1007/978-0-387-87959-8>
- Ma, Q., Liu, J., Liu, Q., Kang, L., Liu, R., Jing, W., ... Liu, M. (2021). Global Percentage of Asymptomatic SARS-CoV-2 Infections Among the Tested Population and Individuals With Confirmed COVID-19 Diagnosis: A Systematic Review and Meta-analysis. *JAMA Network Open*, 4(12), e2137257. <http://doi.org/10.1001/jamanetworkopen.2021.37257>
- Neyman, J. (1937). Outline of a Theory of Statistical Estimation Based on the Classical Theory of Probability. *Philosophical Transactions of the Royal Society of London. Series A, Mathematical and Physical Sciences*, 236(767), 333–380. <http://doi.org/10.1098/rsta.1937.0005>
- Petersen, J. M., Ranker, L. R., Barnard-Mayers, R., MacLehose, R. F., & Fox, M. P. (2021). A systematic review of quantitative bias analysis applied to epidemiological research. *International Journal of Epidemiology*, 50(5), 1708–1730. <http://doi.org/10.1093/ije/dyab061>
- Poole, D., & Raftery, A. E. (2000). Inference for Deterministic Simulation Models: The Bayesian Melding Approach. *Journal of the American Statistical Association*, 95(452), 1244–1255. <http://doi.org/10.1080/01621459.2000.10474324>

- Powers, K. A., Ghani, A. C., Miller, W. C., Hoffman, I. F., Pettifor, A. E., Kamanga, G., ... Cohen, M. S. (2011). The role of acute and early HIV infection in the spread of HIV and implications for transmission prevention strategies in Lilongwe, Malawi: a modelling study. *The Lancet*, 378(9787), 256–268. [http://doi.org/10.1016/S0140-6736\(11\)60842-8](http://doi.org/10.1016/S0140-6736(11)60842-8)
- Reinhart, A., Brooks, L., Jahja, M., Rumack, A., Tang, J., Agrawal, S., ... Tibshirani, R. J. (2021). An open repository of real-time COVID-19 indicators. *Proceedings of the National Academy of Sciences*, 118(51), e2111452118. <http://doi.org/10.1073/pnas.2111452118>
- Robson, B. J. (2014). When do aquatic systems models provide useful predictions, what is changing, and what is next? *Environmental Modelling & Software*, 61, 287–296. <http://doi.org/10.1016/j.envsoft.2014.01.009>
- Rubin, D. B. (1987). The Calculation of Posterior Distributions by Data Augmentation: Comment: A Noniterative Sampling/Importance Resampling Alternative to the Data Augmentation Algorithm for Creating a Few Imputations When Fractions of Missing Information Are Modest: The SIR Algorithm. *Journal of the American Statistical Association*, 82(398), 543. <http://doi.org/10.2307/2289460>
- Rubin, D. B., Gelman, A., & Meng, X.-L. (Eds.). (2004). *Applied Bayesian modeling and causal inference from incomplete-data perspectives: an essential journey with Donald Rubin's statistical family*. Chichester, West Sussex, England ; Hoboken, NJ: Wiley.
- Sah, P., Fitzpatrick, M. C., Zimmer, C. F., Abdollahi, E., Juden-Kelly, L., Moghadas, S. M., ... Galvani, A. P. (2021). Asymptomatic SARS-CoV-2 infection: A systematic review and meta-analysis. *Proceedings of the National Academy of Sciences*, 118(34), e2109229118. <http://doi.org/10.1073/pnas.2109229118>
- Ševčíková, H., Raftery, A. E., & Waddell, P. A. (2007). Assessing uncertainty in urban simulations using Bayesian melding. *Transportation Research Part B: Methodological*, 41(6), 652–669. <http://doi.org/10.1016/j.trb.2006.11.001>